

# Introduction to Couchbase: A NoSQL Document Database

Daniel Owen  
Software Engineer



# Couchbase NoSQL Leadership



## Leading NoSQL database company

Open Source development & business model



## Document-oriented NoSQL database

Focused on interactive internet and mobile applications



## Provide more flexible, higher performance, more scalable database than relational alternative



## Most mature, reliable and widely deployed solution

>7,500 paid production deployments worldwide





## Headquarters in Silicon Valley (Mountain View, CA)

~100 employees including >50 in engineering/product


>80% of commits to Couchbase, memcached, Apache CouchDB

<#>







**Easy Scalability**



**Consistent High Performance**



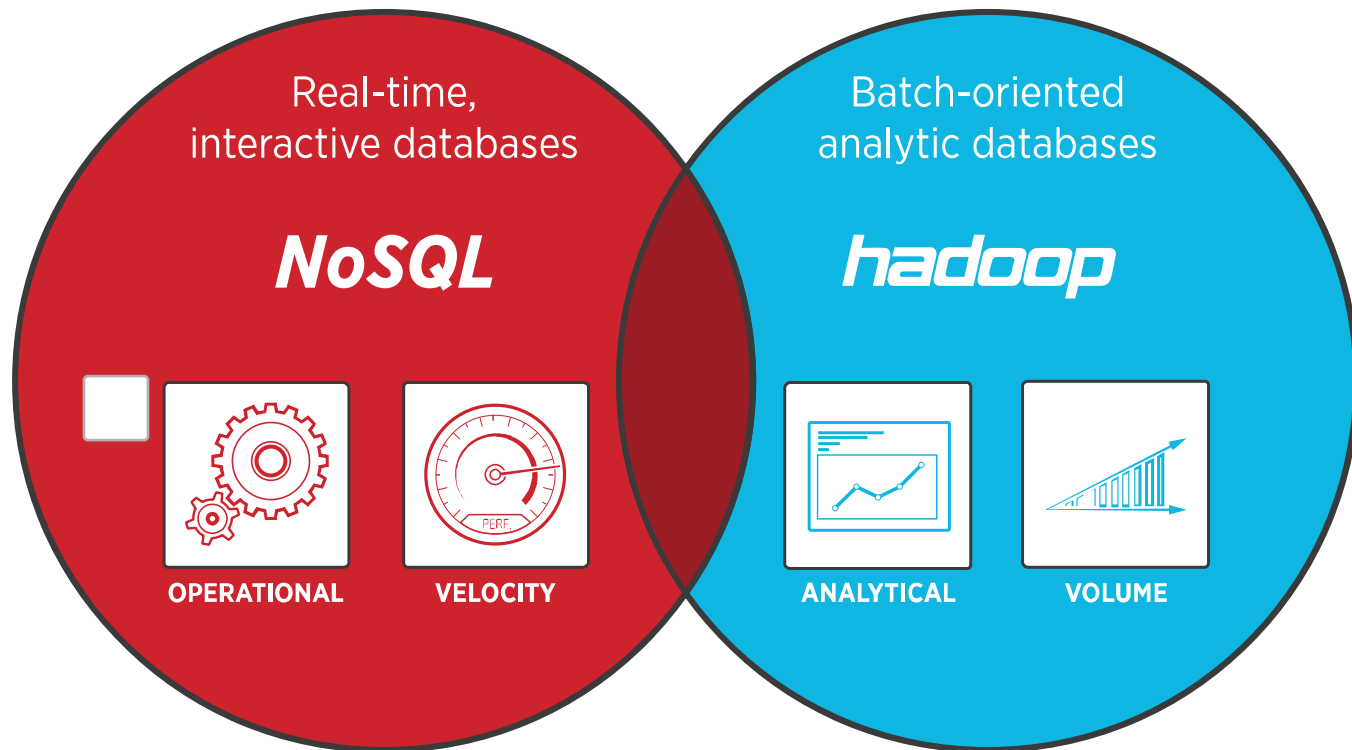
**Always On 24x7x365**



**Flexible Data Model**



# Big Data = Operational + Analytic (NoSQL + Hadoop)



- Online
- Web/Mobile/IoT apps
- Millions of customers/consumers

- Offline
- Analytics apps
- Hundreds of business analysts

# Couchbase Server is Open Source

- All code available on github
- Community Edition for free
- Dedicated to developers
- Try building it yourself!

[github/couchbase/tlm](https://github.com/couchbase/tlm)



# NoSQL == “Not Only SQL”

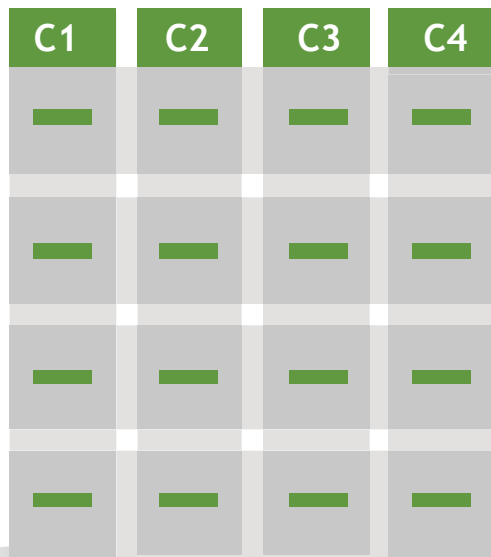
- Users **want** a familiar query language
- SQL has worldwide popularity and recognition
- N1QL allows SQL queries to be run at JSON documents



N1QL = JSON + SQL



# Relational vs Document Data Model



C1	C2	C3	C4
—	—	—	—
—	—	—	—
—	—	—	—
—	—	—	—

## Relational data model

Highly-structured table organization with rigidly-defined data formats and record structure.

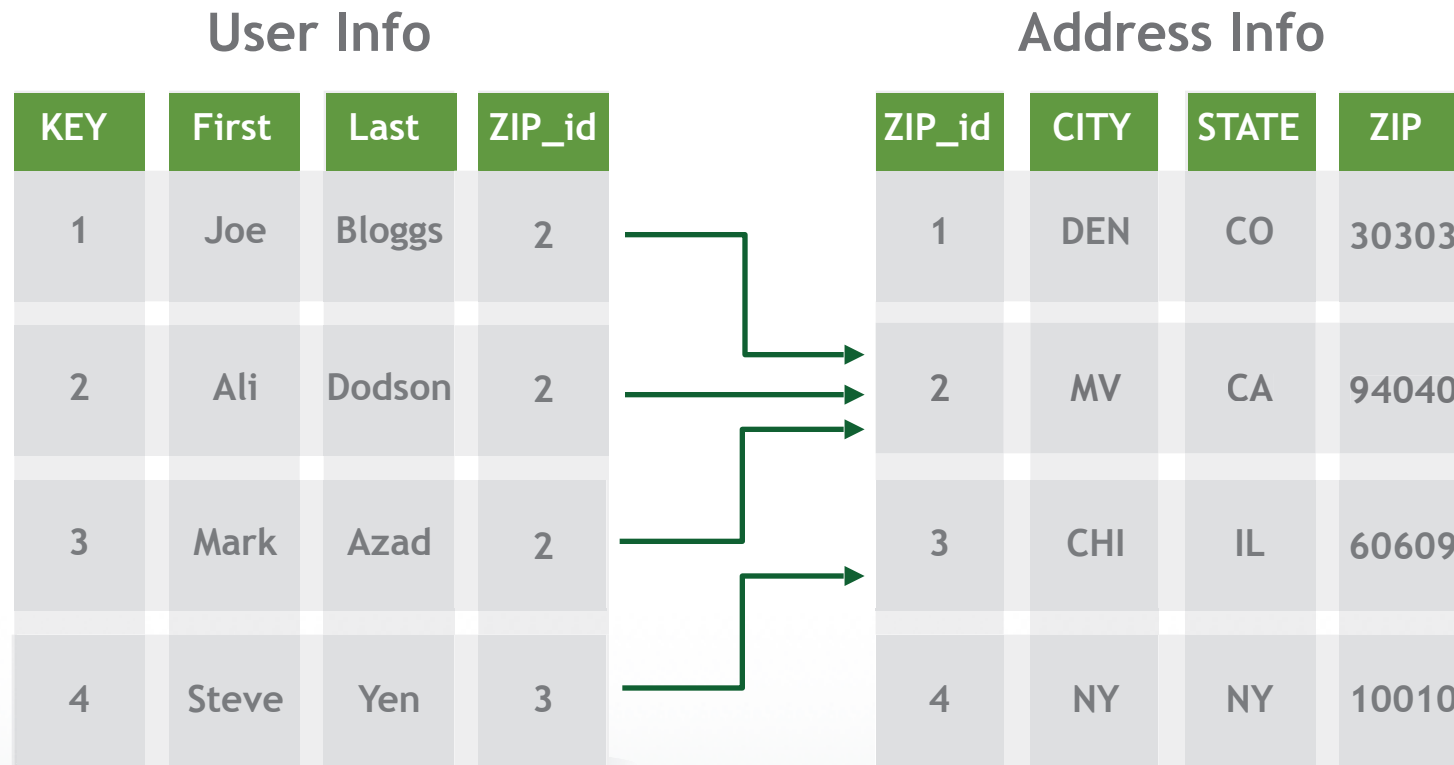


## Document data model

Collection of complex documents with arbitrary, nested data formats and varying “record” format.



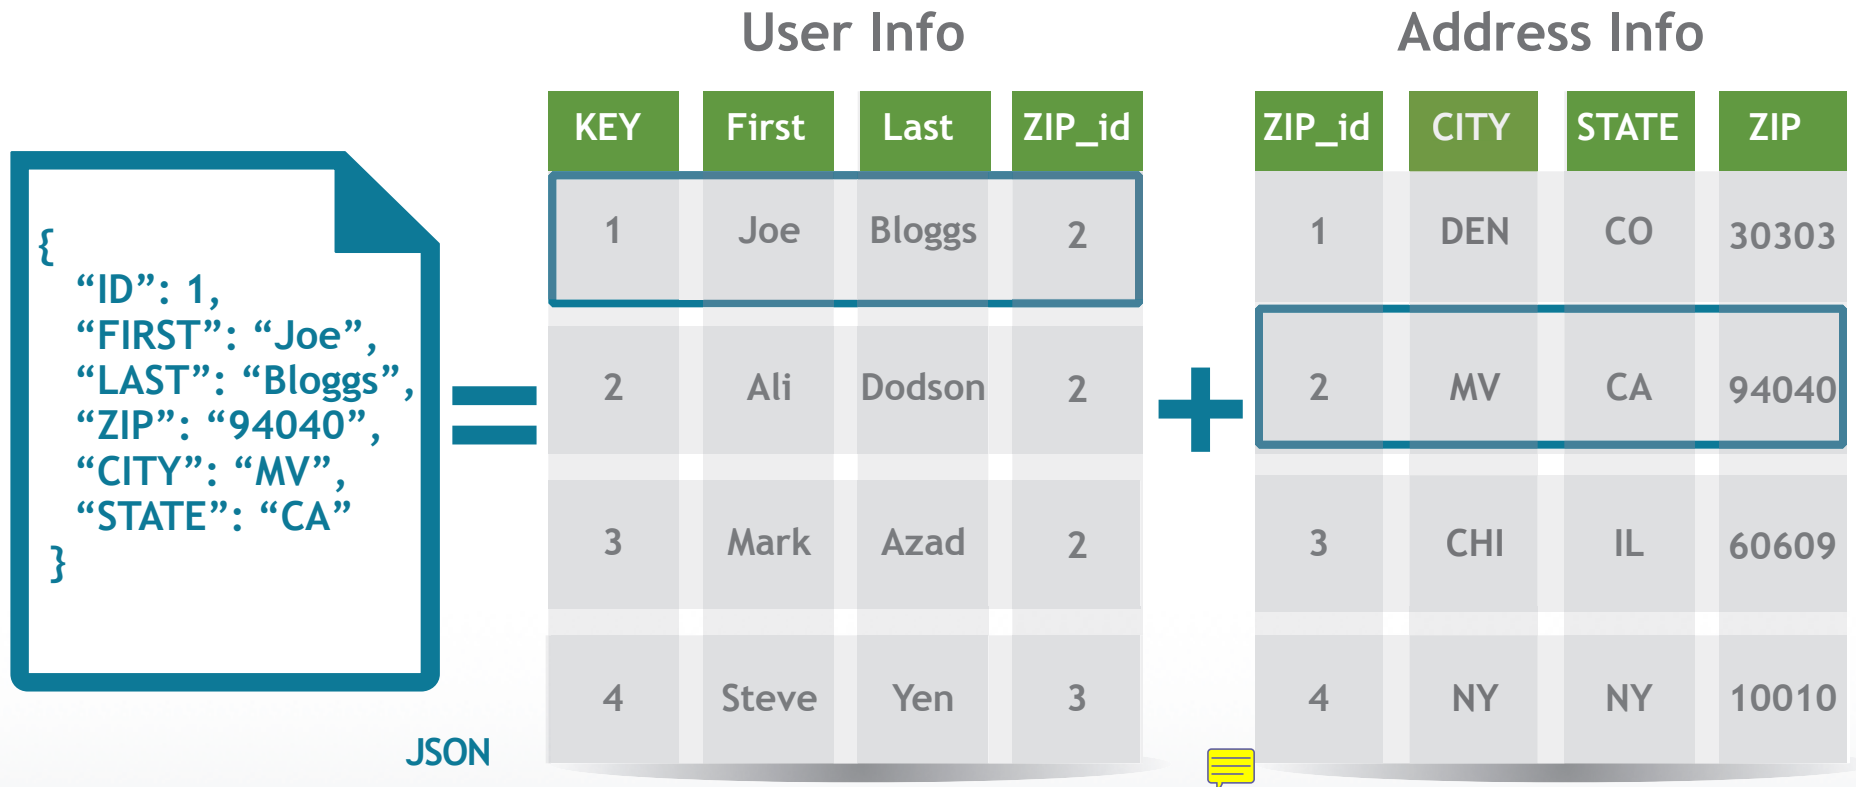
# RDBMS Example: User Profile



To get info about specific user, you perform a join across two tables




# Document Example: User Profile




All data in a single document


# Relational Technology Scales Up



Application Scales Out  
Just add more commodity servers




RDBMS Scales Up  
Get a bigger, more complex server



Expensive and disruptive sharding, doesn't perform at web scale

# Couchbase Server Scales Out Like App Tier



Couchbase Distributed Data Store



Application Scales Out  
Just add more commodity web  
servers

— System Cost  
— Application Performance

Users

NoSQL Database Scales Out  
Cost and performance mirrors app  
tier

— System Cost  
— Application Performance

Users

Scaling out flattens the cost *and* performance curves

# Relational Target Market

Ideal for workloads requiring

## (1) ACID properties

- Atomicity
- Consistency
- Isolation
- Durability

## (2) Ad Hoc Queries \*\*



## (3) High Security

- Full authentication model
- Full auditing
- etc.

\*\* Couchbase currently developing N1QL

## Market

- Transaction Processing
  - Billing
  - ☐ Banking
- High Security
  - Financial market
  - Health industry
- Small number of users
  - Scale-out not required



# NOSQL Market

## Internet Companies

- Social Gaming
- Ad Networks
- Social Networks
- Online Business Services
- E-Commerce
- Online Media
- Content Management
- Cloud Services

## Enterprises

- Communications
- Retail
- Financial Services
- Health Care
- Automotive/Airline
- Agriculture
- Consumer Electronics
- Business Systems

# NOSQL Use Cases

Web app or Use-case	Couchbase Solution
Content Store & Metadata System	Couchbase document store + Elastic Search
Social Game & Mobile App	Couchbase store game and player data
Ad Targeting	Couchbase stores user information for fast access
User Profile Store	Couchbase Server as a key-value store
Session Store	Couchbase Server as a key-value store
High Availability Caching Tier	Couchbase Server as a memcached tier replacement
Chat/Messaging Platform	Couchbase Server

# Some of Couchbase Customers



Couchbase

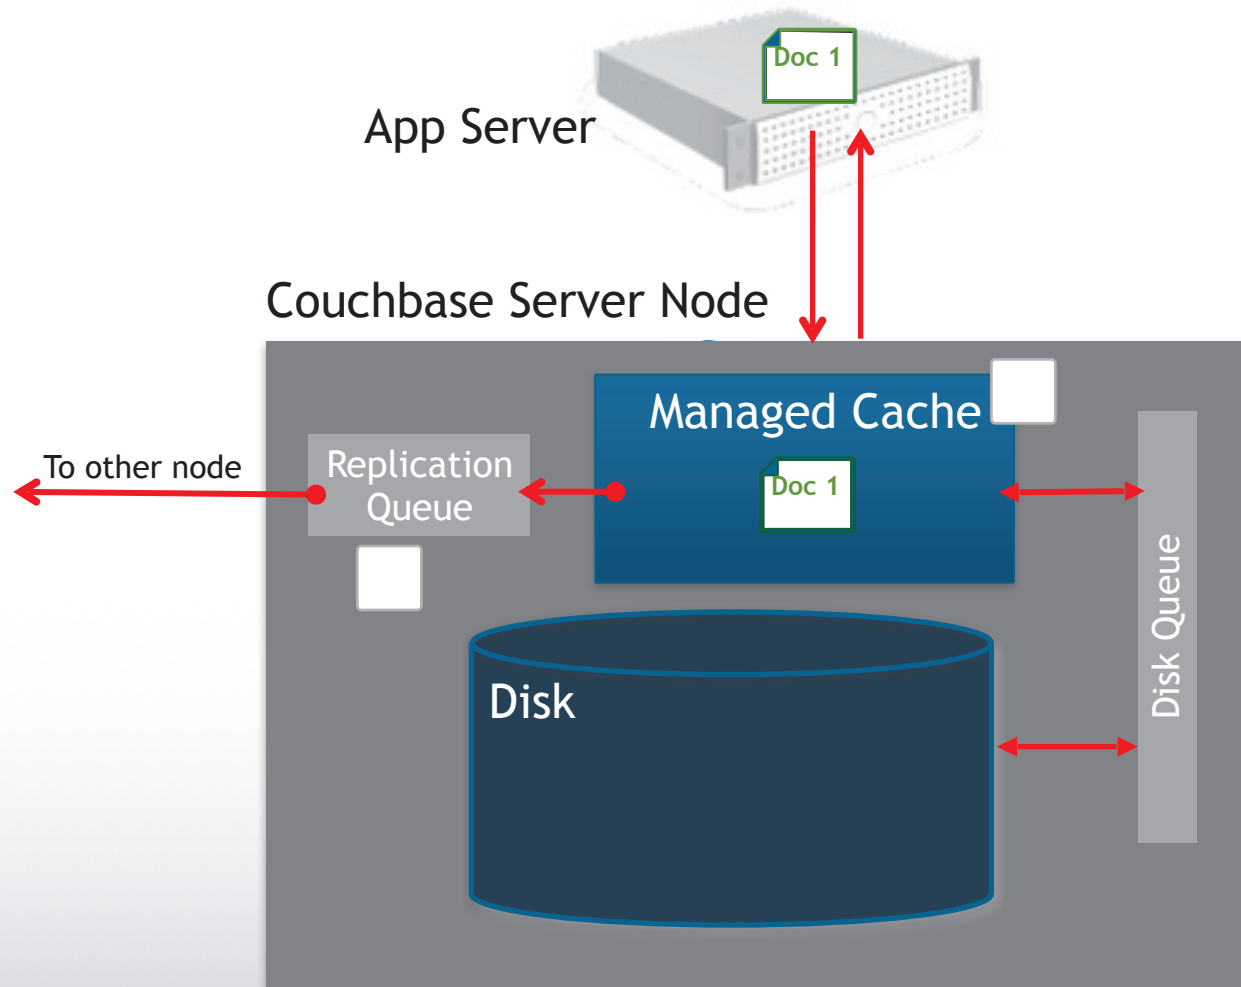




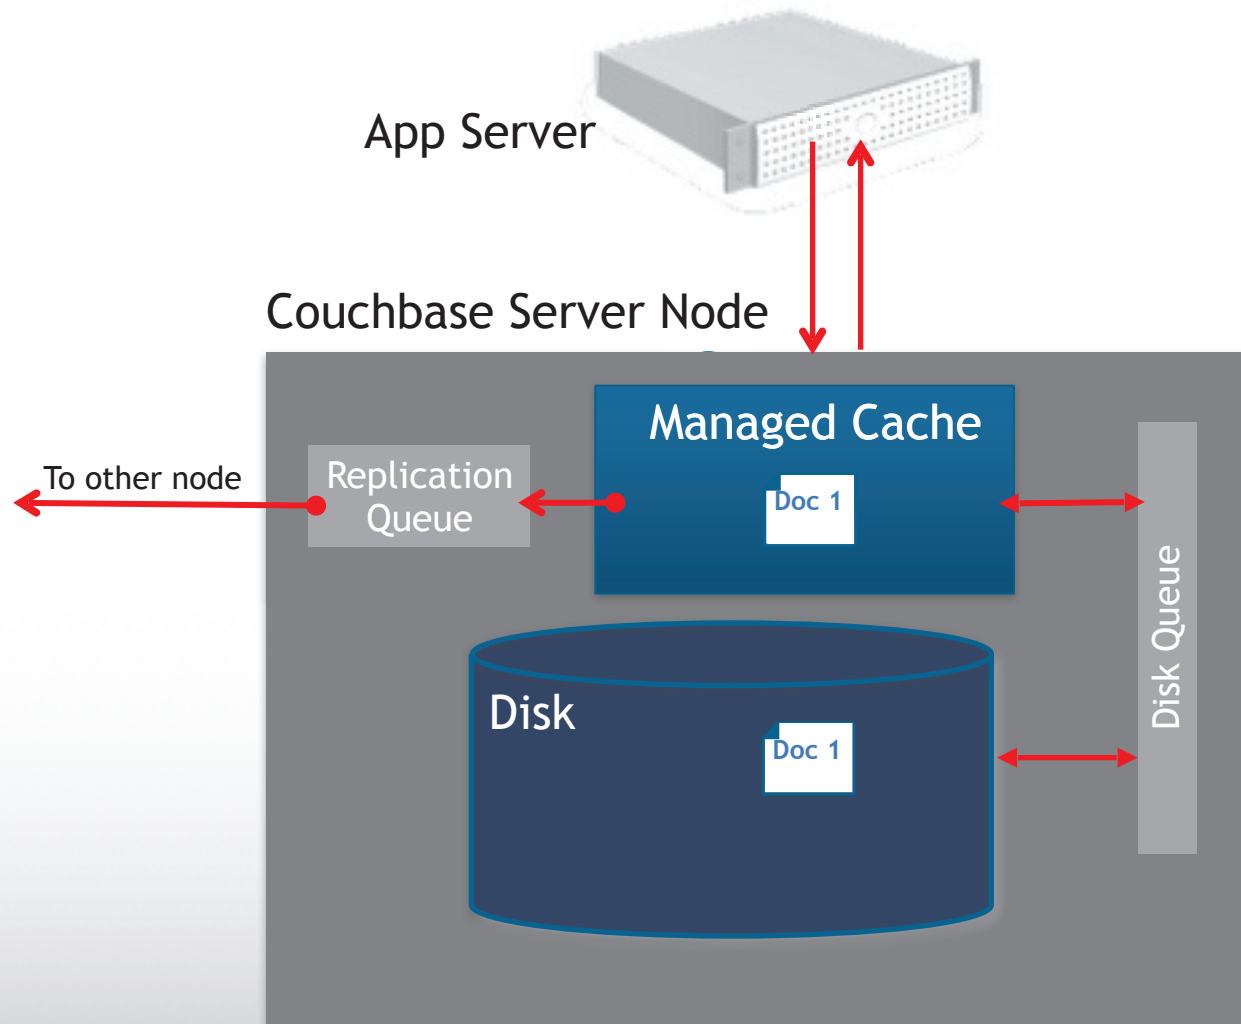
# Couchbase Operations



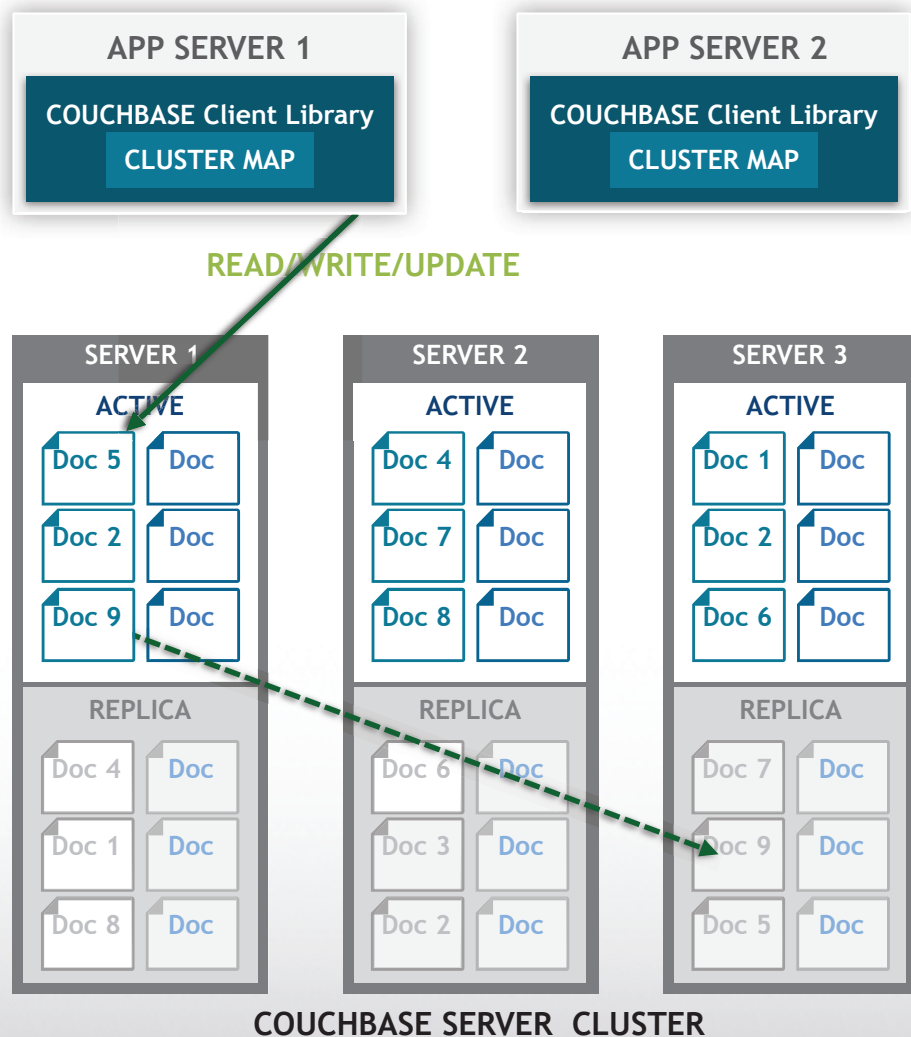
# Single node - Couchbase Write Operation



# Single node - Couchbase Read Operation

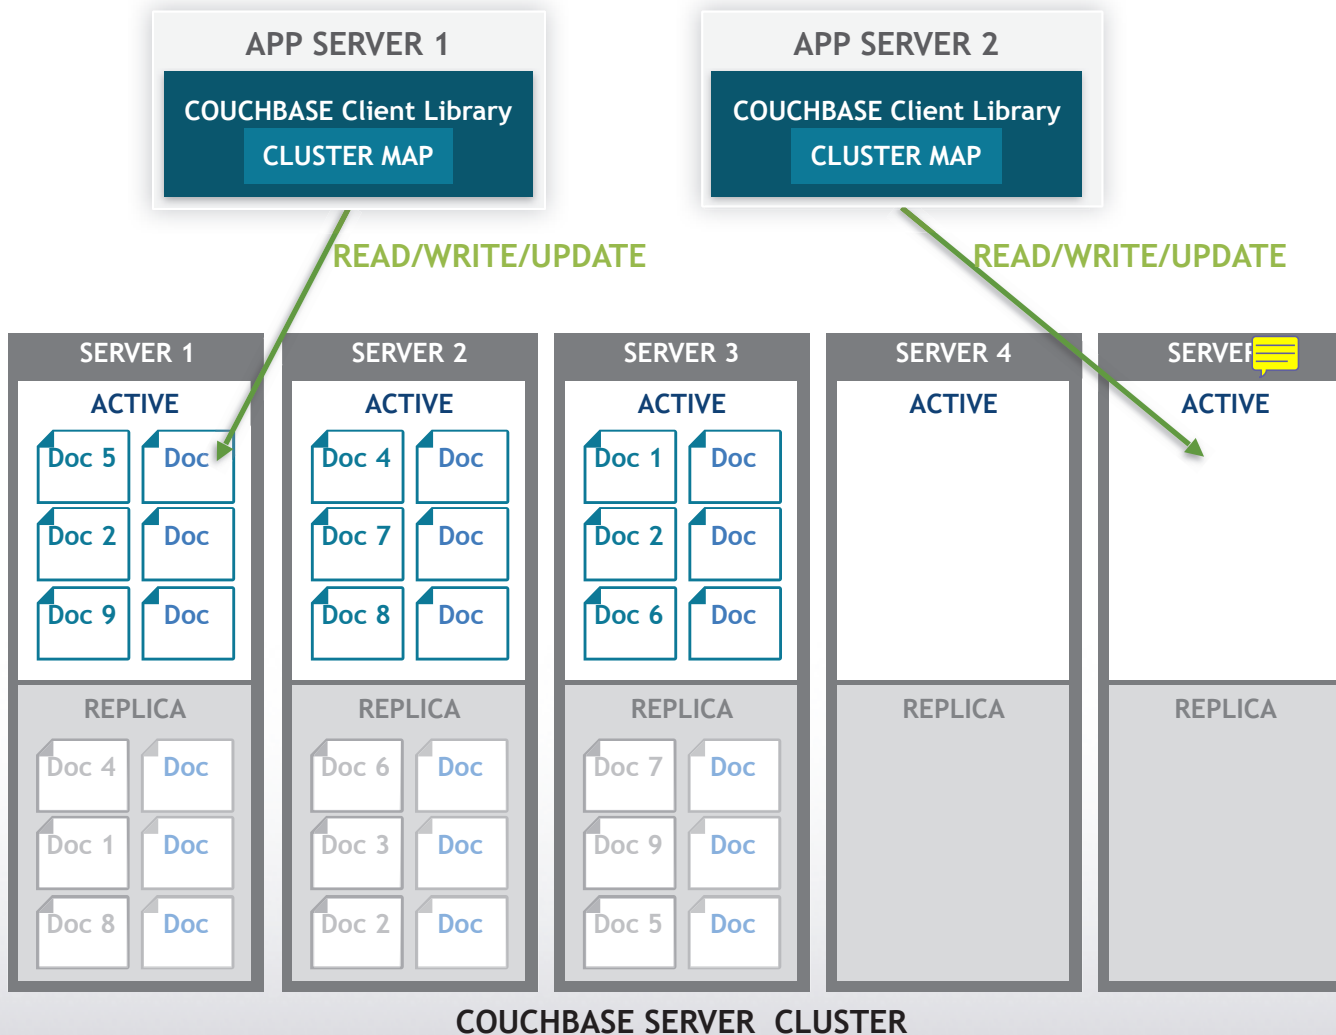


# Basic Operation



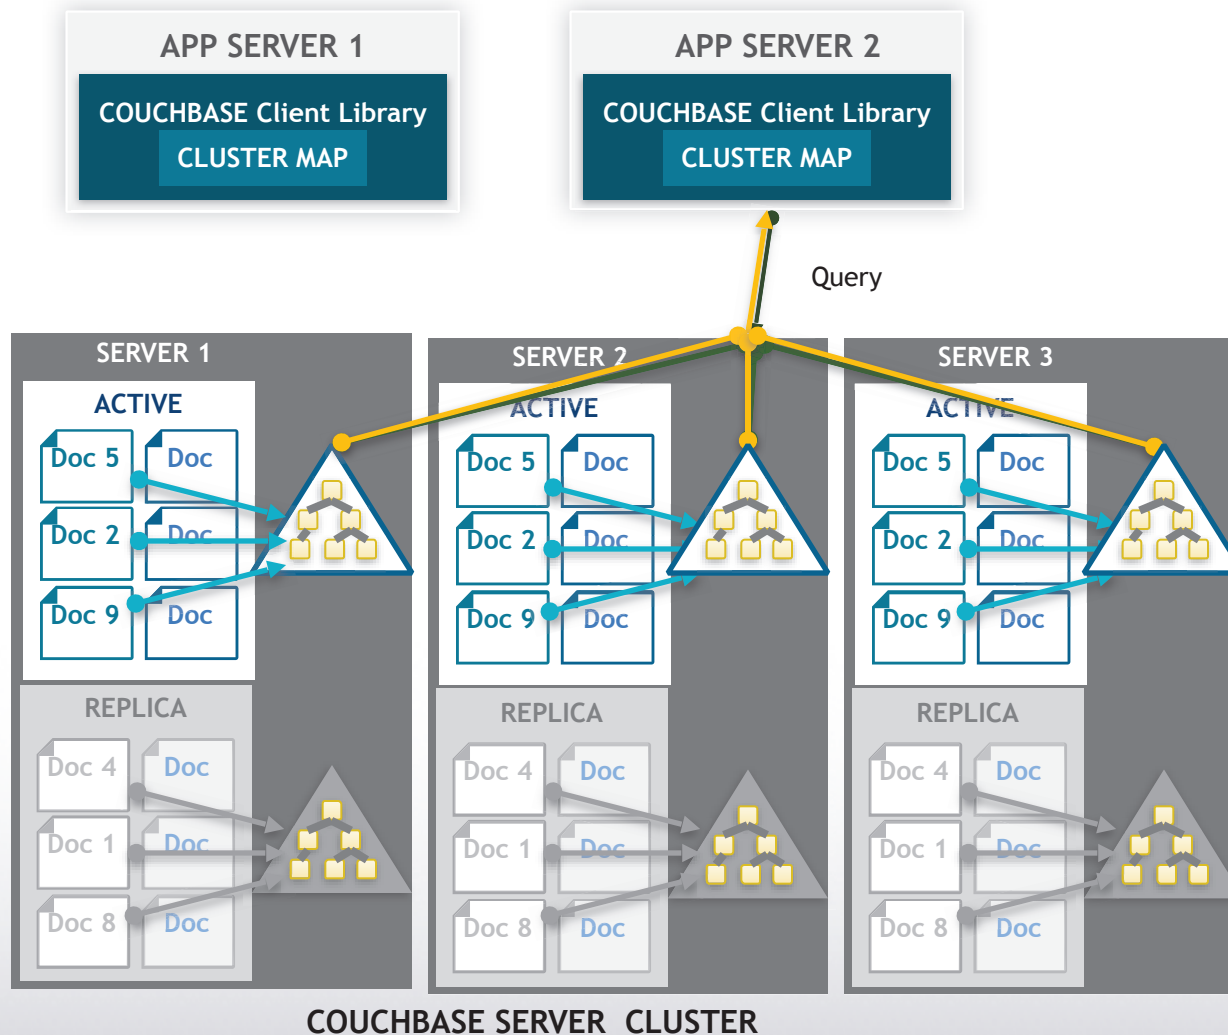
- Docs distributed evenly across servers
- Each server stores both active and replica docs  
Only one server active at a time
- Client library provides app with simple interface to database
- Cluster map provides map to which server doc is on  
App never needs to know
- App reads, writes, updates docs
- Multiple app servers can access same document at same time

# Add Nodes to Cluster



- Two servers added  
One-click operation
- Docs automatically rebalanced across cluster  
Even distribution of docs  
Minimum doc movement
- Cluster map updated
- App database calls now distributed over larger number of servers

# Indexing and Querying



- Indexing work is distributed amongst nodes
- Large data set possible
- Parallelize the effort
- Each node has index for data stored on it
- Queries combine the results from required nodes